



Lecture 4 - Introduction NOSQL, and MongoDB

Instructor: Suresh Melvin Sigera
Email: suresh.sigera@cuny.csi.edu



Agenda

- What is a schema-less aka NoSQL data model?
- CAP theorem
- What is JSON?
- What is MongoDB?
- QA



What is a schema-less data model?

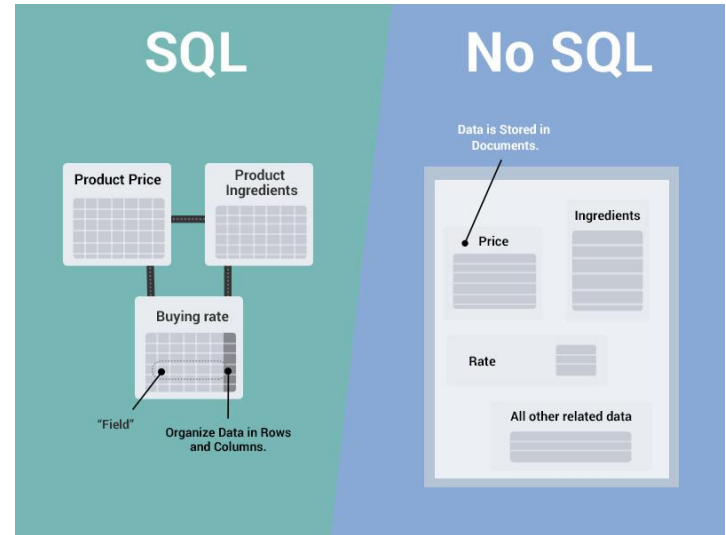
In a relational databases

- You can't add a record which does not fit the schema
- You need to add NULLs to unused items in a row
- We should consider the data type prior to the database design, i.e you can't add a string to an integer field
- You can't add multiple items in a field, you should create another table: primary-key, foreign key, joins and normalization ...!!!

What is a schema-less data model?

In NoSQL databases

- There is no schema to consider
- There is no unused cell
- There is no datatype (implicit)
- Most of the considerations are done in the application layer
- We gather all items as in an aggregate (document)





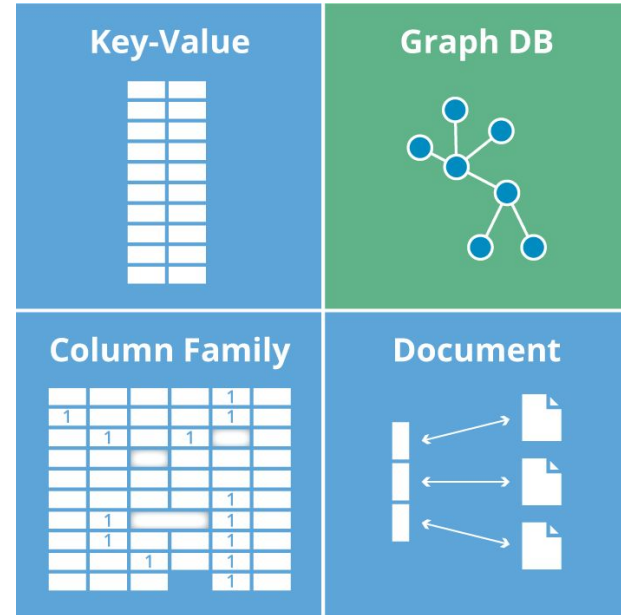
NoSQL databases

What does NoSQL actually mean?

- NoSQL movement = The whole point of **seeking alternatives is that you need to solve a problem that relational databases are a bad fit for**
- NoSQL databases = Next generation databases mostly addressing some of the points: being non-relational, distributed, open-source and horizontally scalable. The original intention has been modern web-scale databases. Often more characteristics apply as: schema-free, easy replication support, simple API, eventually consistent, a huge data amount, and more

Types of NoSQL

- Key-value
- Graph database
- **Document oriented**
- Column family





Benefits of NoSQL

- Scaling
 - Relational databases
 - Traditional approach - scaling up, buying bigger servers as database load increases
 - NoSQL
 - Scale out - distribute data across multiple hosts seamlessly
- Volume AKA Big Data
 - RDBMS
 - Capacity and constraints of data volumes at its limits
 - NoSQL
 - Huge increase in data



Benefits of NoSQL

- Administrators
 - Relational databases
 - Require highly trained expert to monitor DB
 - NoSQL
 - Require less management, automatic repair and simpler data models
- Flexibility
 - Relational databases
 - Change management to schema for RDMS have to be carefully managed
 - NoSQL
 - Databases more relaxed in structure of data
 - Database schema changes do not have to be managed as one complicated change unit
 - Application already written to address an amorphous schema



Benefits of NoSQL

- Economics
 - **Relational databases**
 - **Rely on expensive proprietary servers to manage data**
 - NoSQL
 - Clusters of cheap commodity servers to manage the data and transaction volumes
 - Cost per gigabyte or transaction/second for NoSQL can be lower than the cost for a RDBMS
- Relaxed consistency
 - **Relational databases**
 - **Strong consistency (ACID properties and transactions)**
 - NoSQL
 - Eventual consistency only (BASE properties)
 - I.e. we have to make trade-offs because of the data distribution



The end of relational databases?

- Certainly no
 - RDBMS is a great tool for solving ACID problems
 - When data validity is super important
 - When you need to support dynamic queries
 - NoSQL is great tool for solving data availability problems
 - When it's more important to have fast data than right data; but still it can prevent from application development side
 - When you need to scale bases on the changing requirements
 - Pick the right tool for the job



NoSQL Databases: principles

- Different aspects of data distribution
 - **Scaling**
 - Vertical vs. horizontal
 - Distribution models
 - Sharding
 - Replication: master-slave vs. peer-to-peer architectures
 - CAP properties
 - **Consistency, availability and partition tolerance**
 - ACID vs. **BASE** guarantees



Scalability

What is scalability?


- Capability of a system to handle growing amounts of data and/or queries without losing performance, or its potential to be enlarged in order to accommodate such a growth
- Two general approaches
 - Vertical scaling
 - Horizontal scaling

Scalability

Vertical scaling (scaling up/down)

- Adding resources to a single node in a system
 - E.g. increasing the number of CPUs, extending system memory, using larger disk arrays, ...
 - I.e. larger and more powerful machines are involved
- Traditional choice
 - In favor of strong consistency
 - Easy to implement and deploy
 - No issues caused by data distribution
 - No expensive JOINS

Works well in many cases but ...





Vertical scalability: drawbacks

Performance limits

- Even the most powerful machine has a limit
 - Moreover, everything works well... unless we start approaching such limits
- Higher costs
 - The cost of expansion increases exponentially
 - In particular, it is higher than the sum of costs of equivalent commodity hardware
- Proactive provisioning
 - New projects / applications might evolve rapidly
 - Upfront budget is needed when deploying new machines
 - And so flexibility is seriously suppressed



Vertical scalability: drawbacks


Vendor lock-in

- There are only a few manufacturers of large machines
- Customer is made dependent on a single vendor
 - Their products, services, but also implementation details, proprietary formats, interfaces, ...
- I.e. it is difficult or impossible to switch to another vendor
 - The cost of expansion increases exponentially
 - In particular, it is higher than the sum of costs of equivalent commodity hardware
- Deployment downtime
 - Inevitable downtime is often required when scaling up

Horizontal scalability: fallacies

False assumptions

- Network is reliable
- Latency is zero
- Bandwidth is infinite
- Network is secure
- Topology does not change
- There is one administrator
- Transport cost is zero
- Network is homogeneous




Horizontal scalability: consequences

Significantly increases complexity

- Complexity of management, programming model, ...

Introduces new issues and problems


- Synchronization of nodes
- Data distribution
- Data consistency
- Recovery from failures



Horizontal scalability: conclusion

A standalone node still might be a better option in certain cases.


- E.g. for graph databases - Simply because it is difficult to split and distribute graphs
- In other words
 - **It can make sense to run even a NoSQL database system on a single node**
 - No distribution at all is the most preferred / simple scenario
- But in general, horizontal scaling really opens new possibilities



Horizontal scalability: architecture

What is a cluster?

- = **a collection of mutually interconnected commodity nodes**
- Based on the **shared-nothing architecture**
 - Nodes do not share their CPUs, memory, hard drives, Each node runs its own operating system instance
 - **Nodes send messages to interact with each other**
- Nodes of a cluster can be heterogeneous
- Data, queries, computation, workload, ... this is all **distributed among the nodes** within a cluster





Distribution models



Generic techniques of data distribution

- Sharding
 - **Different data on different nodes**
 - Motivation: increasing volume of data, increasing performance
- Replication
 - Copies of the **same data on different nodes**
 - Motivation: increasing performance, increasing fault tolerance
- Both the techniques are mutually orthogonal
 - I.e. we can use either of them, or combine them both
- Distribution model
 - = specific way how sharding and replication is implemented NoSQL systems often offer automatic sharding and replication

Sharding

Sharding (horizontal partitioning)



- Placement of different data on different nodes
 - What different data means? Different aggregates
 - E.g. key-value pairs, documents, ...
- Related pieces of data that are accessed together should also be kept together
 - Specifically, operations involving data on multiple shards should be avoided



Sharding

Objectives

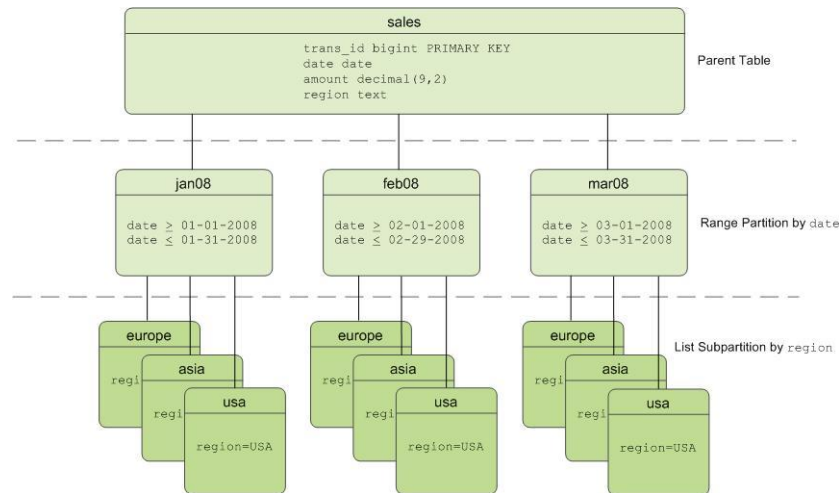
- Uniformly distributed data (volume of data)
- Balanced workload (read and write requests)
- Respecting physical locations
 - E.g. different data centers for users around the world
- Unfortunately, these objectives...
 - may mutually contradict each other
 - may change in time



Sharding

How to actually determine shards for aggregates?

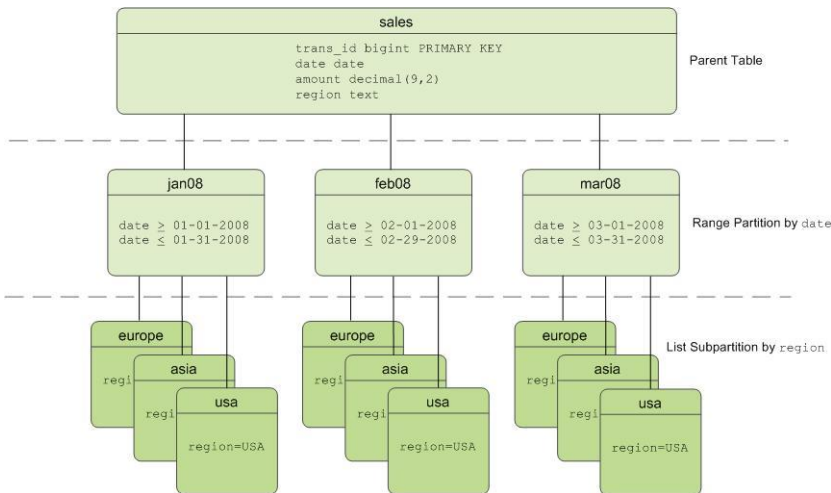
- We not only need to be able to place new data when handling write requests, but also find the data in case of read requests
 - I.e. when a given search criterion is provided (e.g. key, id, ...), we must be able to determine the corresponding shard
 - So that the requested data can be accessed and returned, or failure can be correctly detected when the data is missing



Sharding

Sharding strategies

- Based on mapping structures
 - Placing of data on shards in a random fashion (e.g. round-robin)
 - Mapping of individual aggregates to particular shards must be maintained (this is not a suitable solution)
- Based on general rules: hash partitioning, and range partitioning





Replication

Replication

- Placement of multiple copies – **replicas** – of the **same data** on different nodes
- **Replica on factor** = the number of copies


Two approaches

- Master-slave architecture
- Peer-to-peer architecture

Replication - master-slave architecture

Architecture


- One node is primary (master), all the other secondary (slave)
- Master node bears all the management responsibility
- All the nodes contain identical data



Replication - master-slave architecture


Read requests can be handled by both the master or slaves

- Suitable for read-intensive applications
 - More read requests to deal with → **more slaves to deploy**
 - When the master fails, read operations can still be handled




Replication - master-slave architecture

- Write requests can only be handled by the master
- Newly written replicas are propagated to all the slaves
- Consistency issue
 - Luckily enough, at most one write request is handled at a time
 - But the propagation still takes some time during which obsolete reads might happen
 - Hence certain synchronization on is required to avoid conflicts




Replication - master-slave architecture

- In case of **master failure**, a new one needs to be appointed
 - Manually (user-defined) or automatically (cluster-elected)
 - Since the nodes are identical, appointment can be fast
- Master might therefore represent a **bottleneck**
 - (because of the performance or failures)



Replication - peer-to-peer architecture

- Architecture
 - All the nodes have equal roles and responsibilities
 - All the nodes contain identical data once again
- Both read and write requests can be handled by any node
 - No bottleneck, no single point of failure
 - Both the operations scale well
 - More requests to deal with → more nodes to deploy





The CAP theorem

The CAP theorem, also known as Brewer's theorem, states that **it is impossible for a distributed computer system to simultaneously provide all three of the following guarantees:**

- Consistency (**all nodes** see the same data at the same time)
- Availability (a guarantee that **every request receives a response about whether it was successful or failed**)
- Partition tolerance (the **system continues to operate** despite arbitrary message loss or failure of part of the system)

According to the theorem, a distributed system cannot satisfy all three of these guarantees at the same time.



JSON = JavaScript Object Notation

- Open standard for data interchange; It is a lightweight data-interchange format. It is easy for humans to read and write. It is easy for machines to parse and generate.
- Design goals
 - Simplicity: text-based, easy to read and write
 - Universality: object and array data structures – Supported by majority of modern programming languages
- Derived from JavaScript (but language independent)
- Started in 2002
- File extension: *.json
- Content type: application/json
- <http://www.json.org/>


```
{  
  "title": "Thor",  
  "year": 2018,  
  "actors": [{  
    "firstname": "Chris",  
    "lastname": "Hemsworth"  
  },  
  {  
    "firstname": "Tom",  
    "lastname": "Hiddleston"  
  }  
],  
  "director": {  
    "firstname": "Taika",  
    "lastname": "Waititi"  
  }  
}
```



JSON - data structure

Object

- Unordered collection of name-value pairs (properties)
 - Correspond to structures such as objects, records, structs, dictionaries, hash tables, keyed lists, associative arrays, ...

Example

- `{ "name" : "Tom Cruise", "dob" : 1692 }`
- `{ }`



JSON - data structure

Array

- Ordered collection of values
 - Correspond to structures such as arrays, vectors, lists, sequences ...
 - Values can be of different types, duplicate values are allowed

Example

- `[2, 7, 7, 5]`
- `["Tom Cruise", 1962, 5.7]`
- `[]`



JSON - data structure

Value

- Unicode string
 - Enclosed with double quotes
 - Backslash escaping sequences
 - Example: `"a \n b \" c \\ d"`
- Number
 - Decimals integers or floats
 - Example: `1`, `-0.5`, `1.5e3`
- **Nested Object**
- **Nested Array**
- Boolean value: `true`, `false`
- Missing information: `null`



Keep in mind, JSON is NOT!

- Overly complex
- A "document" format
- A markup language
- A programming language





MongoDB

- **JSON document database**
- Features
 - Open source, high availability, eventual consistency, automatic sharding, master-slave replication, automatic failover, secondary indices
- Developed by MongoDB
- Implemented in C++, C, and JavaScript
- Operating systems: Windows, Linux, Mac OS X ...
- Initial release in 2009



mongoDB®

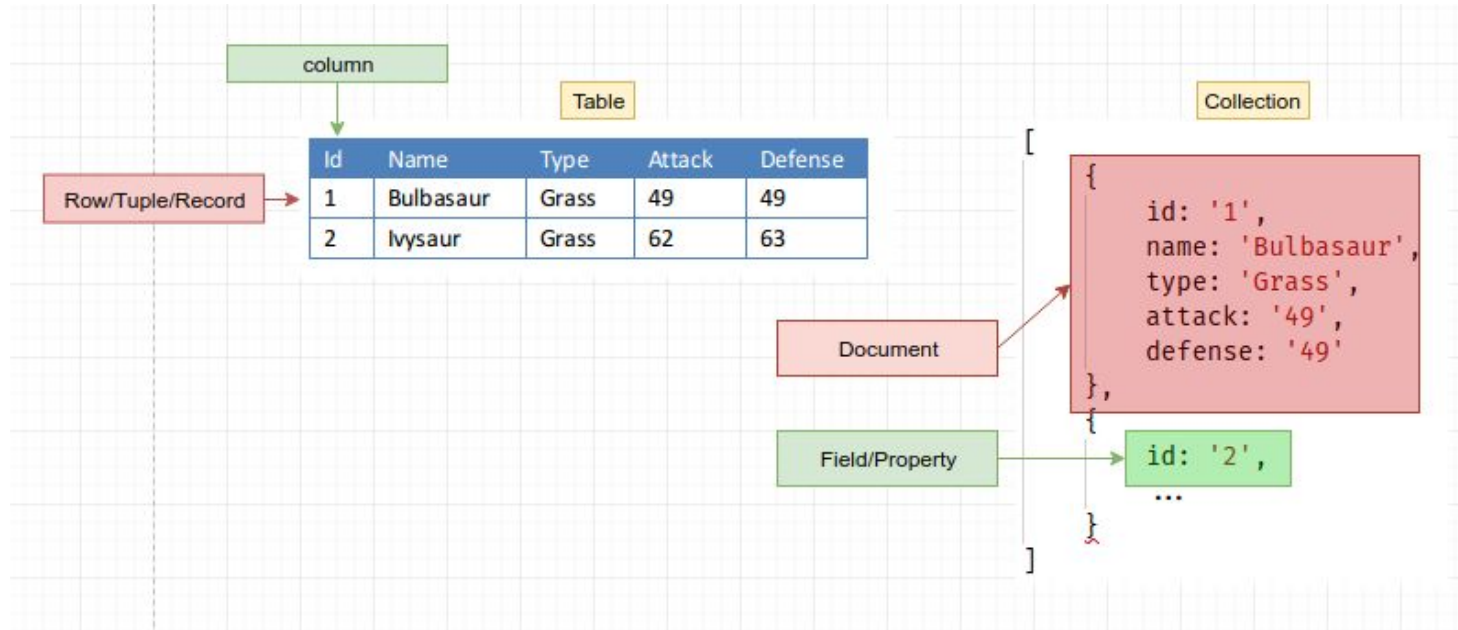


MongoDB - data model

Database system structure: **Instance** → **databases** → **collections** → **documents**

- Database
- Collection
 - Collection of documents, usually of a similar structure
- Document
 - MongoDB document = one JSON object
 - Each document
 - belongs to exactly one collection
 - has a unique identifier `_id`

MongoDB - data model





SQL vs MongoDB concepts

SQL Terms/concepts	MongoDB terms
Database	Database
Table	Collection
Row	Document or BSON Document
Column	Field
Index	Index



SQL vs MongoDB concepts

SQL Terms/concepts	MongoDB terms
Table Joins	Embedded Documents and Linking
Primary Key	Primary Key
Transaction Begin, Commit/Rollback	NA
Schema	NA



MongoDB - data model

MongoDB document = one JSON object

- **Internally stored as BSON (Binary JSON)**
- Maximal allowed size: 16 MB (in BSON)
 - GridFS can be used to split large files into smaller chunks
- Restrictions on field names
 - _id (at the top level) is reserved for a primary key
 - **Field names cannot start with \$**
 - Reserved for query operators
 - **Field names cannot contain .**
 - Used when accessing nested fields



MongoDB - data model

Primary keys

- Features of identifiers
 - **Unique** within a collection
 - **Immutable** (cannot be changed once assigned)
 - Can be of **any type** other than an array
- Design of identifiers
 - Natural identifier
 - Auto-incrementing number – not recommended
 - UUID (Universally unique identifier)
 - **ObjectId - special 12-byte BSON type** (default option)
 - Small, likely unique, fast to generate, ordered, based on a timestamp, machine id, process id, and a process-local counter



JSON-style Documents represented as BSON

```
{ "hello": "world" }
```

```
\x16\x00\x00\x00\x02hello\x00\x06\x00\x00  
\x00world\x00\x00
```



MongoDB - design questions

Flexible schema

- No document schema is provided, nor expected or enforced
 - However, **documents within a collection are similar in practice**
 - **MongoDB document = one JSON object**
 - I.e. even a complex JSON object with other recursively nested objects, arrays or values
- Design challenge
 - Balancing application requirements, performance aspects, and data retrieval patterns,
 - UUID (Universally unique identifier)
 - while considering structure of data and mutual relationships

Two main concepts: references vs. embedded documents



MongoDB - why use JSON?

JSON is open, human and machine-readable standard, and it supports all basic data types: numbers, strings, boolean values, arrays and hashes.

More importantly,

- No need for a fixed schema. Just add what you want
- Format is programmer friendly. Unlike you know... XML!
- MongoDB uses BSON behind the scenes which is an extended format for JSON
- MongoDB reuses a Javascript engine for their queries. Makes complete sense in the world to re-use JSON for object representation



MongoDB - list databases

It's strange to listen but true that MongoDB doesn't provide any command to create databases. Then the question is how would we create database ?. The answer is – We don't create database in MongoDB, we just need to use database with your preferred name and need to save a single record in database to create it. Let's check the current databases in our system.

```
> show dbs
```

```
admin    0.000GB  
local    0.000GB
```



MongoDB - use new database

Now if we want to create database with name CSC424. The keyword **use** will instantiate database object. However, MongoDB doesn't create any database yet, until you save something inside.

```
> use CSC424
```

```
switched to db CSC424
```



MongoDB - create a collection

MongoDB stores the data in the form of JSON documents. A group of such documentation is collectively known as a collection in MongoDB. Thus, a collection is analogous to a table in a relational database while a document is analogous to a record. To store documents, we first need to create a collection. The exciting thing about a NoSQL database is that unlike SQL database, you need not specify the column names or data types in it.

```
> db.createCollection("students")
```

```
{ "ok" : 1 }
```



MongoDB - show collections

If you fancy to see the collections you have it in the current db instance, use the following command.

```
> show collections
```

```
students
```



MongoDB - drop database

MongoDB provides `dropDatabase()` command to drop currently used database with their associated data files. Before deleting make sure what database are you selected using `db` command.

```
> db
```

```
students
```

```
> db.dropDatabase()
```

```
{ "dropped" : "CSC424", "ok" : 1 }
```



QA